CRM Fabric

Specification audit and remediation register

Source: paste-1.md, 'THE CRM FABRIC - Final Complete Build Prompt', 261 lines, 21 KB. Audit date 14 November 2026.
Scope: the document as written. No code exists yet, so this audits the specification, not an implementation.

Bottom line

The document is a build prompt, not an engineering specification. It states intent with high conviction and leaves 43 defects that block or degrade construction, of which 7 are blockers: the builder cannot produce a conforming system from this text without inventing design decisions the document forbids it from asking about.

Three findings carry the weight. First, three of the eight declared layers - Intelligence, Protocol/Settlement, Economic Engine - have no section, no objects and no entry criteria; they exist only as rows in a table. Second, every composite score in the product is undefined: the two routing formulas are unweighted sums of incommensurable terms, trust_score is consumed four times and defined zero times, and RVS is five adjectives. Third, twenty-two deliverables carry no acceptance criteria, so 'done' has no meaning and every build batch self-certifies.

Two findings are more serious than the defect list suggests because they are structural. Four v1 features - Markdown mirror, graph sidecar, RVS, review queue - cannot run in P0 Bare, the profile that is required to pass the full test suite; and three independent history mechanisms are specified with no authority rule, so the product's central question ('what did each system know at time T') has three possible answers.

The commercial claim is also overstated. Of the 40 listed inventions, 18 are documented prior art, 10 are genuinely novel or partly novel, and 10 are speculative or undefined. The differentiation list is roughly half prior art.

The specification is repairable without changing the product thesis. Every defect in this register has a stated fix, and the companion specification (CRM_Fabric_Spec_v1.1.docx) closes all 43 by explicit assumption or formal definition.

> Correction to an earlier reading: the 72-hour breach clock in the specification is not a GDPR transplant. DPDP Rules 2025, rule 7(2), requires detailed intimation to the Data Protection Board within 72 hours of becoming aware. The GDPR transplant is the 30-day DSAR clock, not the breach clock. Defect E-01 is narrowed accordingly.

What the document is, and what it is not

It is a complete statement of product intent: eight layers, forty inventions, a thirty-sprint sequence, an explicit AI contract, and a set of prohibitions that are unusually disciplined. The 'DO NOT' section is the strongest part of the document and should survive any rewrite intact.

It is not a specification. It contains no interface contracts, no error taxonomy, no acceptance criteria, no load model, no data dictionary, no decision log format and no version. Those are the artefacts a builder needs

and the artefacts this register measures the absence of.

Method and evidence base

Every defect cites a line range in the source. Claims about statutory timing were checked against primary sources on 14 November 2026: DPDP Rules 2025 (Gazette notification 13 November 2025) for breach notification and commencement, and the AI Act service desk consolidated text for high-risk application dates following the AI Omnibus. Novelty classifications were made against named prior systems, not against intuition. Performance and schedule figures are labelled estimates wherever they are estimates; nothing in this register is presented as measured.

Defect register

Forty-three defects, ordered by severity then area. Severity means: Blocker - construction cannot proceed correctly; Critical - a stated guarantee has no enforcement point; Major - the artefact is incomplete in a way a reviewer will find; Minor - ambiguity that costs a decision, not a rewrite.

ID	Sev	Location	Defect	Fix
A-01	Blocker	Lines 15-26	Layers 3, 4 and 5 are named in the layer table with target versions but have no section anywhere in the document.	Write entry/exit criteria and object models, or move them to a roadmap section and retitle the document.
A-02	Blocker	Line 4, 261	'Do not ask clarifying questions' is combined with roughly twelve undefined mechanisms: trust_score, RVS, both routing scores, history authority, revocation propagation, autonomy tier distinction.	Replace with an assumption-register protocol: proceed, but record every assumption with a decision ID; stop only at irreversible or security-irreversible forks.
B-01	Blocker	Line 41	Federated routing score is a bare sum of seven incommensurable terms plus two penalties. No weights, no units, no normalization, no tie-break rule.	Fixed weight vector, per-term [0,1] transform, versioned weights logged with every decision.
B-02	Blocker	Lines 33, 41, 45, 81	trust_score is a routing input, an occupant attribute and a conflict tiebreaker, and is never computed. Conflict resolution consumes it; conflicts are its natural input.	Declare formula, observation window, minimum sample, prior and decay.
C-01	Blocker	Lines 56, 149, 154-161	Markdown mirror needs S3, the graph sidecar needs Neo4j, RVS needs the graph. P0 Bare is 'Postgres + app + web' and must pass the full test suite.	Publish a per-profile degradation table and make the test suite profile-aware.
C-02	Blocker	Lines 58, 158	'Every write is a commit' against p95 write <50 ms local ack. A durable commit requires fsync plus object write.	Amortised commit model with a separate commit-latency budget.
G-01	Blocker	Lines 189, 209-232	Twenty-two deliverables, zero acceptance criteria. The build log requires an 'acceptance report' that reports against nothing.	Attach a testable acceptance criterion to each deliverable.
A-03	Critical	Line 19, 70-89	Layer 2 is labelled 'v1 simplified' but its object model is specified at full fidelity with no v1/v2 boundary.	Enumerate the v1 object subset explicitly.
A-04	Critical	Line 87	Only one layer boundary has an interface contract ('assign_to_seat is the only bridge'). Mesh-to-core, core-to-experience and workforce-to-ledger have none.	Define one interface per layer boundary, each with an enforcement test.

ID	Sev	Location	Defect	Fix
B-03	Critical	Line 156	Relationship Velocity Score is described as graph math over frequency, latency asymmetry, meeting density, centrality and decay, with no formula, scale, calibration set or decay constant.	Publish the formula and a calibration fixture.
B-04	Critical	Lines 41, 81	None of the eight routing terms or two penalty terms has a declared range or transform. utilization and cost_efficiency are open-ended.	Declare a transform per term.
B-05	Critical	Lines 58, 66	Latency targets state no concurrency, payload size, hardware class, index state, or cache condition.	Attach a load model to every target.
C-03	Critical	Lines 34, 158, 162	Three independent history mechanisms - bi-temporal columns, git commits, hash-chained events - with no authority rule for 'what each system knew at time T'.	One authority, the others derived. See section 11.
C-04	Critical	Line 62	'Immutable audit log with SHA-256 chain' on append-only Postgres is tamper-evident only against a database owner who is not the attacker. No external anchor, no verification job, no fork or repair procedure.	Anchor the chain head externally on a schedule; publish the verification procedure.
C-05	Critical	Lines 152, 197	P2 bundles the review queue with the MCP container and the local model, while Stage 3 (v1) requires both MCP tools and the review queue.	Move the review queue to P1; keep MCP and models in P2.
D-01	Critical	Line 85	v1 autonomy is 'shadow_only and confirm_then_run only', and 'all agent actions to review queue'. If every action goes to review, the two tiers are behaviourally identical.	Define an observable difference, or collapse to one tier.
D-02	Critical	Line 89	'Removing human stops agent instantly' has no mechanism: no revocation path, no in-flight handling, no propagation deadline, no fail-closed rule.	Signed RevocationEvent, poll interval, in-flight abort policy, fail-closed default.
E-01	Critical	Line 62	DSAR lifecycle is '30-day tracking', a GDPR transplant. DPDP does not set a 30-day DSAR deadline in that form.	Parameterise per jurisdiction; confirm with counsel before shipping as a product claim.
E-02	Critical	Line 136	'AI Act classification per AI action (v2+)'. Transparency obligations are enforceable from 2 Aug 2026; Annex III high-risk rules apply from 2 Dec 2027 following the AI Omnibus.	Classify at v1; defer only high-risk conformity work.
F-01	Critical	Lines 7, 12	Three incompatible ICPs in one v1: solo founder, 3-10 person B2B agency, 5,000-employee organisation.	Name one primary, one secondary, one design constraint.
G-02	Critical	Lines 191-203	Thirty sprints with no team size, velocity or capacity assumption.	State team shape and re-baseline.
A-05	Major	Lines 20, 132-136	Cross-cutting concerns are three bullets, but deliverable 22 requires full specifications of them.	Promote to full sections.
A-06	Major	Line 39	Connector interface lists eight methods with no error taxonomy, pagination contract, partial-failure semantics, or schema-drift detection.	Add error classes and drift detection to the interface.
B-06	Major	Lines 7, 58	'No perceptible slowdown as data grows' is unbounded and contradicts the 1M-record bound two pages later.	Replace with a stated scaling curve and a re-test trigger.
B-07	Major	Line 58	No measurement harness, percentile window or regression gate is defined for any latency target.	Define harness, window and CI gate.

ID	Sev	Location	Defect	Fix
C-06	Major	Lines 56, 148-152	Flutter/SQLCipher mobile is a v1 deliverable and appears in no deployment profile.	Add mobile rows to the profile matrix.
C-07	Major	Line 152	The review queue is a deterministic feature but is bundled into the AI profile, while the spec forbids labelling deterministic features as AI.	Decouple.
D-03	Major	Line 79	Escalation ladders terminate at an unspecified state and can exhaust in a human-only deployment.	Declare terminal states and their SLA effect.
D-04	Major	Line 83	'SLA clock carries across handoffs' with no rule for a receiving seat that declares a different SLA.	Re-evaluate SLA at handoff; record both clocks.
D-05	Major	Line 75	Seats are versioned, but there is no rule for in-flight tasks when permissions_scope or required_capabilities change.	Pin the task to the seat version at assignment.
E-03	Major	Line 136	No DPIA, no record of processing, no sub-processor register, no residency option at v1, for a multi-tenant SaaS fanning contact data across four external nodes.	Add the four artefacts as v1 deliverables.
E-04	Major	Line 62	Consent is modelled as a field (consent_state), not as notice plus purpose-bound consent.	Model notice and purpose, not just state.
F-02	Major	Lines 12, 72	Seat-based workforce value requires multi-occupant organisations. The stated primary ICP has one occupant.	State the solo-founder value path, or change the ICP.
F-03	Major	Line 11	'Aha moment at minute 4' has no event, no funnel step and no definition of the end state.	Name the activation event.
G-03	Major	Line 193	Stage 1 builds the HubSpot connector before the canonical schema is versioned and before conflict resolution exists.	Version the schema and land conflict resolution before the second connector.
G-04	Major	Line 203	External security audit sits in Stage 6 with no pre-audit hardening gate and no remediation sprint.	Add a hardening gate and reserve two sprints.
G-06	Major	Lines 177-185	Forty inventions claimed with no novelty audit. At least twenty are documented prior art.	Reclassify every entry. See section 8.
H-01	Major	Whole d ocument	Eleven wedge-and-moat triad constructions, nine one-line closers, fourteen subjectless imperatives, six unbounded absolutes ('never', 'instantly', 'always', 'no perceptible').	Rewrite absolutes as bounded claims; cut the triads.
D-06	Minor	Lines 77-78	Task.handoff_chain and HandoffPacket have no declared relationship.	Declare the type.
E-05	Minor	Lines 209-232	No LICENSE, SECURITY.md, DPA template or EULA among the 22 deliverables.	Add to Stage 5.
F-04	Minor	Lines 11, 116, 256	Three clocks - 4-minute aha, 10-minute onboarding, 30-minute switching cost - with no stated relationship.	State which clock is the gate.
G-05	Minor	Line 189	Build-log fields are named with no format, schema or retention rule.	Publish a schema.
H-02	Minor	Whole d ocument	No version, author, date, status or change log on the document itself.	Add document control.

Counts by severity: Blocker 7, Critical 14, Major 17, Minor 5. By area: architecture 6, metrics 7, runtime 7, workforce 6, compliance 5, product 4, plan 6, document 2.

Three of eight layers are placeholders

The layer table names eight layers. Five have bodies. Layers 3, 4 and 5 appear once, in the table, and nowhere else. The document jumps from Layer 2 to Layer 6 with no note that three sections are missing - a reader following the sequence will not notice the gap.

Layer	Name	Declared	Lines	Status	Finding
0	Federation Mesh	v1	28-47	Specified	Object model complete; routing score unimplementable as written; ledger authority undefined.
1	CRM Core	v1 (optional)	49-68	Specified	Stack and performance stated; four features cannot run in P0; history mechanisms conflict.
2	Workforce	v1 simplified, v2 full	70-89	Partial	Full-fidelity model with no v1 boundary; autonomy tiers indistinguishable in v1.
3	Intelligence	v2	-	Absent	Named in the layer table only. No section, no objects, no entry criteria.
4	Protocol / Settlement	v3	-	Absent	Named in the layer table only.
5	Economic Engine	v3	-	Absent	Named in the layer table only.
6	Experience	v1	91-120	Specified	Strongest section in the document. Absolute 'no spinners' rule conflicts with asynchronous reads.
7	Learning	v1	122-130	Specified	Ten artefacts listed; no acceptance criterion for any of them.

Fix: either write the three layers, or retitle the document 'five layers, three roadmap items' and move them to the roadmap section with entry criteria. Leaving them in the layer table while absent from the body is the one defect that damages the document's credibility on first read.

Every composite score is undefined

Seven scores drive routing, trust, forecasting, conflict resolution and relationship health. None has a formula that two engineers would implement identically. This is the single largest source of hidden rework: the router is specified as deterministic, and determinism is exactly the property that is missing.

Score	Terms	Missing	Proposed definition
Federated routing score	7 additive terms, 1 penalty	Weights, units, normalization, tie-break	score = sum(w_i * t_i) - 0.20 * cross_boundary_penalty, every t_i in [0,1]. Weights versioned and logged per decision.
Workforce routing score	6 additive terms, 1 penalty	Weights, units, normalization, escalation_penalty definition	Same form as federated score; escalation_penalty = 0.3 if the candidate is itself a fallback seat, else 0.
trust_score	Referenced four times, defined zero times	Formula, window, sample floor, prior, decay	EWMA over 90 days of (success, 1-escalation, 1-conflict, audit completeness); n>=20; prior 0.5; clamped [0,1].
Relationship Velocity Score	5 qualitative drivers	Formula, scale, decay constant, calibration set	RVS = sum_i decay(dt_i) * w_type(type_i) * centrality(entity) / (1 + latency_asymmetry), decay half-life 21 days, normalized to tenant p95.
Forecast (weighted pipeline)	'weighted' with no weighting scheme	Stage weights, override rule, provenance link	Stage probability from trailing 12-month conversion per stage, recomputed nightly, each number linked to its evidence tree.

Score	Terms	Missing	Proposed definition
cost_efficiency	Named as a term	Cost model, budget source	1 - (estimated_cost / seat_cost_budget), clamped [0,1]; estimated_cost from node cost_model.
data_freshness	Named as a term	Staleness function, half-life per node	exp(-age_hours / half_life_hours), half-life 24 h for sales nodes, 4 h for support nodes.

The failure mode is not that the terms are wrong. capability_match, sla_headroom and cost_efficiency are the right terms. The failure is that an unweighted sum lets utilization dominate a router that is supposed to be capability-first, and nothing in the document says otherwise. Fix: a fixed weight vector, a declared [0,1] transform per term, and a version stamp on every logged decision.

Four v1 features have no runtime

The AI contract requires P0 Bare - Postgres, app, web - to pass the full test suite. Four v1 features require infrastructure P0 does not have. Either P0 fails its own gate or four v1 features are silently optional. Neither is stated.

Feature	P0	P1	P2	P3	Conflict
Postgres + app + web	yes	yes	yes	yes	None
NATS event bus	no	yes	yes	yes	None - CQRS log lives in Postgres at P0
Read replica	no	yes	yes	yes	None
Redis	no	yes	yes	yes	None
Neo4j graph sidecar + RVS	no	yes	yes	yes	C-01: RVS is a v1 feature and is absent at P0
S3 Markdown mirror	no	no	no	no	C-01: v1 feature with no profile that runs it
Review queue	no	no	yes	yes	C-05: required by Stage 3 (v1) but only present at P2
MCP container (10 tools)	no	no	yes	yes	C-05: required by Stage 3 (v1)
Local model (Ollama / llama.cpp)	no	no	optional	optional	None
Agent passports, autonomy tiers	no	no	no	yes	Consistent with v1 autonomy = shadow + confirm
Flutter / SQLCipher mobile	not stated	not stated	not stated	not stated	C-06: v1 deliverable with no profile
i18n (5 locales)	yes	yes	yes	yes	None
External audit anchoring	not stated	not stated	not stated	not stated	C-04: required by the immutability claim

The review-queue row is the sharpest instance: the queue is not an AI feature, yet it is bundled into the AI profile and simultaneously required by Stage 3. Moving it to P1 resolves both the profile conflict and the AI-contract contradiction in one edit.

History has three authorities and no ruling

The document specifies bi-temporal columns on every entity, git commits on every write, and a hash-chained UniversalEvent log. Each is a complete history mechanism. Together they are three answers to the same question, and the document never says which one is true.

This matters more than it looks. 'What each system knew at time T' is the product's central query - it feeds the digital twin, the dispute trail, provenance and the audit chain. Three mechanisms that can disagree means three possible answers, and a customer dispute will find the disagreement.

Resolution: the ledger is authoritative; bi-temporal columns are a projection of it; git is a checkpoint log over the ledger, not a per-write commit stream. That single ruling also removes the 50 ms-write versus fsync-per-commit contradiction, because commits become batched rather than per-record.

Half the invention list is prior art

Forty inventions are claimed. Classified against named prior systems: 18 prior art, 10 partly novel or novel framing, 10 speculative, undefined or narrative, 2 research-stage or contested. The novelty claim is not a marketing problem; it is a build-order problem, because prior art should be adopted and novel work should be built, and the document treats all forty the same way.

No	Invention	Ver	Class	Basis
1	Bi-temporal records	v1	Prior art	Snodgrass 1994; implemented in XTDB, Datomic, SQL:2011 system-versioned tables.
2	Identity spine + federated relationship graph	v1	Partly novel	Stable UUIDv7 + per-cell provenance + graph overlay is a distinct combination.
3	Local-first real-time as one sync problem	v1	Prior art	Local-first literature; Replicache, ElectricSQL, Automerge.
4	ZK consent proofs	v2	Research	No production precedent for DPDP consent at the stated timeframe.
5	Edge-sync with server-offloaded validation	v1	Prior art	Offline-first plus transactional outbox is a standard pattern.
6	Conversation-to-record	v1	Prior art	Email and ICS ingestion is table stakes in every CRM.
7	Autonomy-aware review queue	v1	Partly novel	Approval queues are standard; binding them to seat autonomy tiers is not.
8	Portable record contracts	v1	Partly novel	Data portability as a contractual artefact with signed attestation.
9	Seat-based versioned job descriptions	v1	Novel framing	RBAC plus capacity plus versioning, framed as an org design object.
10	Mixed-fleet queues	v1	Prior art	Workforce management plus agent orchestration.
11	Cost-aware routing	v2	Prior art	Task routing and auction literature.
12	Escalation ladder as first-class object	v1	Prior art	PagerDuty and ServiceNow escalation policies.
13	Handoff packets with full state and signature	v1	Partly novel	Full-state handoff with signature is distinct; the concept is not.
14	Shadow mode / apprenticeship	v1	Prior art	Shadow deployment and canary release.
15	AI shift scheduling	v1	Prior art	Workforce management.
16	Playbook authoring with deterministic compilation	v1	Partly novel	The AI-proposal-to-deterministic-workflow gate is distinct.

No	Invention	Ver	Class	Basis
17	Workforce analytics with provenance	v2	Prior art	Data lineage.
18	Agent-to-agent commerce	v2	Contested	Active standardisation; A2A and MCP ecosystems. Not a moat.
19	Customer-side agent bridge	v2	Speculative	No schema, no counterparty, no adoption path.
20	Cross-company relationship graph	v2	Partly novel	Requires bilateral data sharing that no v1 customer has agreed to.
21	Marketplace with provenance + revenue share	v2	Prior art	App marketplaces.
22	Company digital twin	v2	Undefined	Term used without a model.
23	Outcome-based billing	v2	Prior art	Usage and outcome billing exist in metering infrastructure.
24	Self-improving org	v2	Speculative	Proposals only, no mechanism.
25	Protocol-based interop	v2	Prior art	MCP, A2A, OpenAPI.
26	Sovereign deployment	v4	Prior art	Self-hosted and air-gapped deployment.
27	Dream Layer	v2	Undefined	Named once, never described.
28	Universal org chart protocol	v2	Speculative	No schema given.
29	Reputation economy for seats	v4	Partly novel	Trust score as a portable asset.
30	Sovereign pipeline	v4	Narrative	A positioning claim, not an invention.
31	Federated CRM mesh	v1	Partly novel	Data virtualisation exists; the ledger plus authority-map combination is distinct.
32	Cross-company chains	v3	Speculative	No counterparty model.
33	Universal event ledger	v1	Prior art	Event sourcing, audit logs, Kafka.
34	Capability-aware work router	v1	Prior art	Capability-based routing is standard.
35	Zero-migration adoption	v1	Prior art	Federation and virtualisation.
36	Cross-CRM conflict resolution with provenance	v1	Partly novel	The combination with a logged authority map is distinct.
37	Chain-of-chains	v3	Undefined	Named, never described.
38	Fair-share distribution	v3	Prior art	Fair queuing.
39	Universal dispute trail	v3	Speculative	No adjudication model.
40	Mesh-level digital twin	v3	Speculative	No model.

Twelve entries are the ones worth defending: 2, 7, 8, 9, 13, 16, 20, 29, 31, 36, plus 3 and 5 as well-executed prior art. The rest should be labelled honestly. A defensible claim of twelve distinct mechanisms is stronger than an indefensible claim of forty, and it survives diligence.

Thirty sprints, no capacity assumption

The plan covers 22 deliverables across 30 two-week sprints - roughly 14 months - with no team size, velocity or capacity stated anywhere. Scope includes Django, four connectors, a local-first sync engine, Flutter with SQLCipher, Neo4j, NATS, an MCP surface, 80 icons, a component library, WCAG 2.2 AA, five locales, a three-system migration wizard, five videos, a docs site, a community, an external audit and a ten-customer beta.

Estimate, with method stated: 22 deliverables at a median of two sprints each is 44 sprints before parallelisation; at a six-engineer, one-designer, half-SRE baseline with realistic overlap, 44 is the planning number and 30 is the aspiration. The estimate is arithmetic on the deliverable list, not measured velocity. It should be replaced with measured velocity after Stage 1.

Two ordering risks are independent of capacity. Stage 1 builds the HubSpot connector before the canonical schema is versioned, guaranteeing schema churn under a live connector. Stage 6 places an external security audit after launch assets are built and budgets no remediation sprint, so audit findings arrive when they cannot be fixed before launch.

The compliance substrate is GDPR-shaped

The substrate is written against European defaults: 30-day DSAR, 72-hour breach, consent as a state field. Checked against current Indian and EU law, one of those three holds, one does not, and one is incomplete.

Holds: the 72-hour clock. DPDP Rules 2025, rule 7(2), requires detailed intimation to the Data Protection Board within 72 hours of awareness, with immediate first intimation and notice to each affected Data Principal without delay. The specification's breach workflow is aligned.

Does not hold: the 30-day DSAR clock. It is a GDPR transplant with no equivalent statutory basis asserted in the document. Parameterise it per jurisdiction and do not publish it as a statutory commitment until counsel confirms.

Incomplete: consent. DPDP consent is notice plus purpose-bound, specific, informed and unambiguous consent, with a Consent Manager regime for registered intermediaries. A single consent_state field cannot carry purpose limitation. Model notice and purpose as objects.

AI Act timing, checked against the consolidated text: the Act applies from 2 August 2026; transparency obligations are enforceable from that date; Annex III high-risk obligations were deferred by the AI Omnibus to 2 December 2027, with Annex I products at 2 August 2028. Deferring 'AI Act classification per AI action' to v2 therefore leaves v1 agent occupants unclassified during a live transparency-obligation period. Classify at v1; defer only the high-risk conformity work.

Also absent from all 22 deliverables: DPIA, record of processing, sub-processor register, LICENSE, SECURITY.md, DPA template, EULA. Enterprise procurement blocks on the last four.

Prose patterns that hide untestable claims

The writing is dense and confident, and that confidence is doing work the specification should do. Five patterns recur, and each one converts a testable requirement into a rhetorical one.

Pattern	Count	Example and consequence
Wedge-and-moat triads	11	'CRM is the wedge. The mesh is the moat. The company operating system is the destination.' Rhythm substitutes for argument. Cut to one claim.
One-line closers	9	'Move.' 'Ship v1.' 'Record v2-v4.' Instruction-shaped fragments that carry no testable content.
Subjectless imperatives	14	'Do not ask clarifying questions.' 'Do not silently overwrite.' No actor, so no owner and no enforcement point.
Unbounded absolutes	6	'instantly', 'never', 'always', 'no perceptible', 'zero-migration', 'no spinners'. Each becomes an untestable acceptance criterion.
Numeric claims without basis	4	'1M records', '50 ms', '100 ms', 'minute 4'. No load model, no measurement method, no owner.

The fix is mechanical: every absolute becomes a bounded claim with a number and a test, every subjectless imperative gets an actor, and every triad is cut to its strongest clause. The document loses about 8% of its length and gains a testable surface.

Remediation order

The defects are not independent. Five decisions unlock the rest, and making them in this order avoids rebuilding artefacts twice.

#	Action	Closes	Blocks
1	Name one primary ICP; demote the other two to secondary and design constraint.	F-01, F-02	Every prioritisation decision
2	Rule on history authority: ledger authoritative, bi-temporal projected, git as checkpoints.	C-02, C-03	RVS, twin, disputes, audit
3	Publish the metric dictionary: weights, transforms, trust, RVS, freshness.	B-01 to B-04	Router, conflicts, forecast
4	Reconcile deployment profiles and move the review queue to P1.	C-01, C-05, C-07	P0 gate, Stage 3 scope
5	Attach acceptance criteria to all 22 deliverables.	G-01	Every acceptance report
6	Write the assumption register and permit recorded assumptions in place of questions.	A-02	All remaining gaps
7	Reclassify the 40 inventions and re-baseline the schedule.	G-02, G-06	Roadmap, diligence
8	Correct the compliance substrate to DPDP and classify AI actions at v1.	E-01 to E-04	Enterprise sale
9	Write or strike Layers 3-5; write cross-cutting specs.	A-01, A-05	Document credibility
10	Prose pass: bound the absolutes, cut the triads.	H-01	Testability of everything above

Steps 1 to 5 change the document's meaning and must precede any build. Steps 6 to 10 can run in parallel with Stage 0.

What survives unchanged

The product thesis, the mesh-first sequencing, the AI contract and the DO NOT list are sound and should not be renegotiated in the rewrite. Specifically: no AI in the critical path; AI disabled by default; AI writes proposals not commits; no agent seat without a fallback; no escalation without a full-state handoff; no silent conflict overwrite; no chain dependent on one node. Each of those is an enforceable rule, and each is rarer in a v1 specification than it should be.

The Experience layer is the strongest section in the document: the jargon map is specific, the preserve-versus-reinvent split is a real product decision, and the lint rule that enforces it is the kind of mechanism most specifications omit. The Learning layer is the second strongest. Both should survive the rewrite with only acceptance criteria added.